

L05 Configuration Management

Markus Raab

Institute of Information Systems Engineering, TU Wien

21.04.2021

This work is licensed under a Creative Commons “Attribution-ShareAlike 4.0 International” license.



History of CM

- 1 History of CM
- 2 CM Languages
- 3 CM Tools
- 4 Meeting

Learning Outcomes

Students will be able to

- remember differences between CM languages and historical approaches.
- write simple configuration management scripts.

Definition

Configuration Management:

- is a discipline in which configuration (in the broader sense) is administered.
- makes sure computers are assembled from desired parts and the correct applications are installed.
- ensures that the execution environment of installed applications is as required.

Definition

Configuration Management Tools:

Definition

Configuration Management Tools:

- help people involved in configuration management.
- have means to describe the desired configuration of the whole managed system.
- try to converge the actual configuration to the desired one [4].

Challenging tasks in configuration management:

Challenging tasks in configuration management:

- inventory list
- installing packages
- monitoring
- add/replace machines
- maintaining files/databases
- ***configuration file manipulation***

Cloning

It all started with:

- clone all files with dd, rdist, rsync or unison (“golden image”)
- then do necessary modifications with scripts or profiles

Cloning

It all started with:

- clone all files with dd, rdist, rsync or unison (“golden image”)
- then do necessary modifications with scripts or profiles
 - + works good for many identical stateless machines
 - fails if differences between machines are too big

Scripts

First improvement: have a script to create the “golden image”.

Possible benefits:

- Documentation

Scripts

First improvement: have a script to create the “golden image”.

Possible benefits:

- Documentation
- Customization (using configuration settings)

Scripts

First improvement: have a script to create the “golden image”.

Possible benefits:

- Documentation
- Customization (using configuration settings)
- **Reproducibility**: Reproduce creation using different operating system versions

Profiles

Profiles are groups of configuration settings between which the user can easily switch.

- by hostname, information EEPROM, manual selection, ...
- can be activated via the profile plugin:

```
1 [application/profile]
2   type := string
3   opt := p
4   opt/long := profile
5   default := current
```

with a config like:

```
1 application/current/key = "current"
2 application/myprofile/key = "myprofile"
3 application/%/key = "default"
```

First four configuration management tools

Cloning, and then NIS/NFS, was state of the art for a long time, until in 1994 when *“the community nearly exploded with four new configuration systems”* [6]:

lcfg from Anderson [2]. The development of lcfg started first in 1991 [1, 2]. Nevertheless, its development still continues [3, 9].

GeNUAdmin from Harlander [7].

omniconf from Hideyo [8].

config from Rouillard and Martin [12].

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility
- Declarative description of the system
(Infrastructure as Code [10])

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility
- Declarative description of the system
(Infrastructure as Code [10])
- Less configuration drift

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility
- Declarative description of the system
(Infrastructure as Code [10])
- Less configuration drift
- Error handling

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility
- Declarative description of the system
(Infrastructure as Code [10])
- Less configuration drift
- Error handling
- Pull/Push

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility
- Declarative description of the system
(Infrastructure as Code [10])
- Less configuration drift
- Error handling
- Pull/Push
- **Reusability**

Possible Benefits

- All advantages scripts have:
Documentation, Customization, Reproducibility
- Declarative description of the system
(Infrastructure as Code [10])
- Less configuration drift
- Error handling
- Pull/Push
- Reusability
- (Resource) Abstractions

L05 Configuration Management

Markus Raab

Institute of Information Systems Engineering, TU Wien

21.04.2021

This work is licensed under a Creative Commons “Attribution-ShareAlike 4.0 International” license.



CM Languages

- 1 History of CM
- 2 CM Languages**
- 3 CM Tools
- 4 Meeting

See Reading Text

See accompanied reading text for this section.

L05 Configuration Management

Markus Raab

Institute of Information Systems Engineering, TU Wien

21.04.2021

This work is licensed under a Creative Commons “Attribution-ShareAlike 4.0 International” license.



CM Tools

- 1 History of CM
- 2 CM Languages
- 3 CM Tools**
- 4 Meeting

List of CM tools

- CFengine (1993)

List of CM tools

- CFengine (1993)
- LCFG (1994)

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)

List of CM tools

- CFengine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)
- Chef (2009)

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)
- Chef (2009)
- Salt (2011)

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)
- Chef (2009)
- Salt (2011)
- Ansible (2012)

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)
- Chef (2009)
- Salt (2011)
- Ansible (2012)
- Itamae (2014)

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)
- Chef (2009)
- Salt (2011)
- Ansible (2012)
- Itamae (2014)
- Puppet

List of CM tools

- CFEngine (1993)
- LCFG (1994)
- Quattor (2005)
- Puppet (2005), Bolt (2018)
- Chef (2009)
- Salt (2011)
- Ansible (2012)
- Itamae (2014)
- Puppet
- OpsMops (2019)

Key/value access in puppet-libelektra [11]:

```
1 kdbkey {'/slapd/threads/listener':
2     ensure => 'present',
3     value  => '4'
4     check => {
5         'type' => 'short',
6         'range' => '1,2,4,8,16',
7         'default' => '1'
8     }
9 }
```

Key/value access in puppet-libelektra:

```
1 kdbmount { 'system:/sw/samba':
2     ensure => 'present',
3     file   => '/etc/samba/smb.conf',
4     plugins => 'ini'
5 }
6 kdbkey { 'system:/sw/samba/global/workgroup':
7     ensure => 'present',
8     value  => 'MY_WORKGROUP'
9 }
10 kdbkey { 'system:/sw/samba/global/log_level':
11     ensure => 'absent'
12 }
```

Uniqueness of keys is essential. Ideally, applications already mount their configuration at installation.

Key/value specifications in puppet-libelektra:

```
1 kdbkey { 'system:/sw/samba/global/log level':
2     ensure => 'present',
3     value  => 'MY_WORKGROUP',
4     check => {
5         'type' => 'short',
6         'range' => '0-10',
7         'default' => '1',
8         'description' => 'Sets the amount of log/
9             debug messages that are sent to the
10            log file. 0 is none, 3 is consider-
11            able.'
12 }
```


Key/value specifications in puppet-libelektra:

```
1 kdbkey { 'spec:/xfce/pointers/Mouse/RightHanded':  
2     ensure => 'present',  
3     check => {  
4         'namespaces/#0' => 'user',  
5         'namespaces/#1' => 'system',  
6         'visibility' => 'important',  
7         'default' => 'false',  
8         'check/type' => 'boolean'  
9     } }
```

Ideally, applications already specify their settings.

Key/value access in Chef:

```
1 kdbset 'system:/sw/samba/global/workgroup' do
2     value 'MY_WORKGROUP'
3     action :create
4 end
```

Key/value access in Chef:

```
1 kdbset '/slapd/threads/listener' do
2     value '4'
3     action :create
4 end
```

Key/value access in Chef:

```
1 kdbset '/slapd/threads/listener' do
2     value '4'
3     action :create
4 end
```

Finding

We have CM code representing the settings.

Key/value access in Ansible:

```
1 - name: setup LDAP
2   connection: local
3   hosts: localhost
4   tasks:
5     - name: set listening threads
6       elektra:
7         key: '/slapd/threads/listener'
8         value: '4'
```

Key/value access in Ansible:

```
1 - name: setup samba
2   connection: local
3   hosts: localhost
4   tasks:
5     - name: set workgroup
6       elektra:
7         mountpoint: system:/sw/samba
8         file: /etc/samba/smb.conf
9         plugins: ini
10      elektra:
11        key: 'system:/sw/samba/global/workgroup'
12        value: 'MY_WORKGROUP'
```

Key/Values Revisited

Decide about **changeability** per key:

- Who is responsible (end user, packages, admin manual or CM).

Key/Values Revisited

Decide about **changeability** per key:

- Who is responsible (end user, packages, admin manual or CM).
- In which namespaces apps search the key (cascading lookup).

Key/Values Revisited

Decide about **changeability** per key:

- Who is responsible (end user, packages, admin manual or CM).
- In which namespaces apps search the key (cascading lookup).
- Who can see it (visibility).

Key/Values Revisited

Decide about **changeability** per key:

- Who is responsible (end user, packages, admin manual or CM).
- In which namespaces apps search the key (cascading lookup).
- Who can see it (visibility).
- Who can edit it (admin, end user, both).

Key/Values Revisited

Decide about **changeability** per key:

- Who is responsible (end user, packages, admin manual or CM).
- In which namespaces apps search the key (cascading lookup).
- Who can see it (visibility).
- Who can edit it (admin, end user, both).
- Which configuration values are allowed (validation).

Key/Values Revisited

Decide about **changeability** per key:

- Who is responsible (end user, packages, admin manual or CM).
- In which namespaces apps search the key (cascading lookup).
- Who can see it (visibility).
- Who can edit it (admin, end user, both).
- Which configuration values are allowed (validation).

Changeability

Ownership of every key must be very clear and documented.

Layers of Abstractions

Recursively define useful abstractions (meta-levels):

- Bits in (configuration) files and memory

Layers of Abstractions

Recursively define useful abstractions (meta-levels):

- Bits in (configuration) files and memory
- Key/value view of configuration settings

Layers of Abstractions

Recursively define useful abstractions (meta-levels):

- Bits in (configuration) files and memory
- Key/value view of configuration settings
- Goals/specifications of settings per node and instantiations of modules

Layers of Abstractions

Recursively define useful abstractions (meta-levels):

- Bits in (configuration) files and memory
- Key/value view of configuration settings
- Goals/specifications of settings per node and instantiations of modules
- CM code to instantiate settings in the whole network

Layers of Abstractions

Recursively define useful abstractions (meta-levels):

- Bits in (configuration) files and memory
- Key/value view of configuration settings
- Goals/specifications of settings per node and instantiations of modules
- CM code to instantiate settings in the whole network
- **Global optimization: allocation of nodes and decision regarding topology in the whole network**

Layers of Abstractions

Recursively define useful abstractions (meta-levels):

- Bits in (configuration) files and memory
- Key/value view of configuration settings
- Goals/specifications of settings per node and instantiations of modules
- CM code to instantiate settings in the whole network
- Global optimization: allocation of nodes and decision regarding topology in the whole network
- **Global goals/specifications of the whole network**

Design Rules [5]

- Factor processes into containers to avoid overlaps in settings.

Design Rules [5]

- Factor processes into containers to avoid overlaps in settings.
- Maintain clear separation of ownership (for every key).

Design Rules [5]

- Factor processes into containers to avoid overlaps in settings.
- Maintain clear separation of ownership (for every key).
- Specify replicated settings in a single source (use links and derivations).

Design Rules [5]

- Factor processes into containers to avoid overlaps in settings.
- Maintain clear separation of ownership (for every key).
- Specify replicated settings in a single source (use links and derivations).
- Document all remaining overlaps (in the specification).

Design Rules [5]

- Factor processes into containers to avoid overlaps in settings.
- Maintain clear separation of ownership (for every key).
- Specify replicated settings in a single source (use links and derivations).
- Document all remaining overlaps (in the specification).
- The manageability of settings is reduced by the number of possible configuration values.

Open Topics

- global optimizations/self-healing

Open Topics

- global optimizations/self-healing
- configuration integration

Open Topics

- global optimizations/self-healing
- configuration integration
- safe migrations of settings and data

Open Topics

- global optimizations/self-healing
- configuration integration
- safe migrations of settings and data
- collaboration

Open Topics

- global optimizations/self-healing
- configuration integration
- safe migrations of settings and data
- collaboration
- management (including knowledge)

Open Topics

- global optimizations/self-healing
- configuration integration
- safe migrations of settings and data
- collaboration
- management (including knowledge)
- centralized vs. distributed

Conclusion

- have unique identifier for your configurations settings
→ allows to get/set configurations and specifications

Conclusion

- have unique identifier for your configurations settings
→ allows to get/set configurations and specifications
- solving CM is solving constraints
→ be aware of the specifications

Conclusion

- have unique identifier for your configurations settings
→ allows to get/set configurations and specifications
- solving CM is solving constraints
→ be aware of the specifications
- do not design around tools but design tools around you

Conclusion

- have unique identifier for your configurations settings
→ allows to get/set configurations and specifications
- solving CM is solving constraints
→ be aware of the specifications
- do not design around tools but design tools around you
- be brave and remove all configuration settings you can

Conclusion

- have unique identifier for your configurations settings
→ allows to get/set configurations and specifications
- solving CM is solving constraints
→ be aware of the specifications
- do not design around tools but design tools around you
- be brave and remove all configuration settings you can
- use all help you can get: e.g. build tools, preseeding, installer automation, virtualization, package managers, distributions

Conclusion

- have unique identifier for your configurations settings
→ allows to get/set configurations and specifications
- solving CM is solving constraints
→ be aware of the specifications
- do not design around tools but design tools around you
- be brave and remove all configuration settings you can
- use all help you can get: e.g. build tools, preseeding, installer automation, virtualization, package managers, distributions
- complexity in CM vs. complexity in applications' specification

L05 Configuration Management

Markus Raab

Institute of Information Systems Engineering, TU Wien

21.04.2021

This work is licensed under a Creative Commons “Attribution-ShareAlike 4.0 International” license.



Meeting

- 1 History of CM
- 2 CM Languages
- 3 CM Tools
- 4 Meeting**

- [1] Paul Anderson. Local system configuration for syssies. Technical report, CS-TN-38, Department of Computer Science, University of Edinburgh, Edinburgh, 1991.
- [2] Paul Anderson. Towards a high-level machine configuration system. In *LISA*, volume 94, pages 19–26, 1994.
- [3] Paul Anderson, Alastair Scobie, et al. Lcfg: The next generation. In *UKUUG Winter conference*, pages 4–7, 2002.
- [4] Mark Burgess. A site configuration engine. In *USENIX Computing systems*, volume 8, pages 309–337, 1995.
- [5] Mark Burgess and Alva L Couch. Modeling next generation configuration management tools. In *LISA*, pages 131–147, 2006.
- [6] Rémy Evard. An analysis of UNIX system configuration. In *LISA*, volume 97, pages 179–194, 1997.

- [7] Magnus Harlander. Central system administration in a heterogeneous Unix environment: GeNUAdmin. In *LISA VIII Proceedings*, 1994.
- [8] Imazu Hideyo. OMNICONF—making os upgrades and disk crash recovery easier. In *LISA VIII Proceedings*, 1994.
- [9] Johannes Hintsch, Carsten Göring, and Klaus Turowski. A review of the literature on configuration management tools. 2016.
- [10] Waldemar Hummer, Florian Rosenberg, Fábio Oliveira, and Tamar Eilam. Testing idempotence for infrastructure as code. In David Eysers and Karsten Schwan, editors, *Middleware 2013*, pages 368–388, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg. ISBN 978-3-642-45065-5.

- [11] Markus Raab, Bernhard Denner, Stefan Hahnenberg, and Jürgen Cito. Unified configuration setting access in configuration management systems. In *ICPC '20: 28th International Conference on Program Comprehension, Seoul, Republic of Korea, July 13-15, 2020*, pages 331–341. ACM, 2020. doi: 10.1145/3387904.3389257. URL <https://doi.org/10.1145/3387904.3389257>.
- [12] John P. Rouillard and Richard B. Martin. Config: A mechanism for installing and tracking system configurations. In *LISA*, 1994.